

Implementing Filter Lock and Peterson-based Tree Lock

Kushagra Indurkha
(CS19B1017)

Introduction	2
Filter Lock	3
Tree-based Peterson Lock	4
Peterson Lock Node	4
Constructor	4
Lock and Unlock	4
Binary Tree of locks	6
Overview	6
Constructor	7
Lock and Unlock	8
Analysis	9
References:	10

Introduction

For n thread locking solution I have tried to implement it in two ways:

- Filter Lock: A direct generalization Of Peterson's algorithm for n threads (The Art of Multiprocessor Programming Maurice Herlihy, Nir Shavit 2.4), is implemented exactly as per discussions in class and in the book.
- Peterson Tree Lock: (The Art of Multiprocessor Programming Maurice Herlihy, Nir Shavit 2.10 Exercise 13)
 - **n(number of threads) needs to be in the power of 2**
 - Arrange a number of 2-thread Peterson locks in a binary tree.
 - Each thread is assigned a leaf lock which it shares with one other thread.
 - Each lock treats one thread as thread 0 and the other as thread 1.
 - In the tree-lock's acquire method, the thread acquires every two-thread Peterson lock from that thread's leaf to the root.
 - The tree-lock's release method for the tree-lock unlocks each of the 2-thread Peterson locks that thread has acquired.

For testing these solutions as mentioned in the assignment instructions I create n threads which in turn have critical sections which they call k times.

Both the filter lock and the tree-based lock implements lock and unlock functions which are virtual functions from a lock-base class

```
/* This is the base class for all locks. */  
class lock_base  
{  
public:  
    virtual void lock(int thread_id) = 0;  
    virtual void unlock(int thread_id) = 0;  
};
```

Filter Lock

```
class filter : public lock_base
{
private:
    /* data */
    int *level;
    int *victim;
    int n;

public:
    /* This is the constructor of the filter class. It initializes the level and victim
    arrays to 0. */
    filter(int n)
    {
        level = new int[n];
        victim = new int[n];
        this->n = n;
        for (int i = 0; i < n; i++)
            level[i] = 0;
    }
    /* This is the implementation of the Filter Lock algorithm. */
    void lock(int thread_id)
    {
        int me = thread_id;
        for (int i = 1; i < n; i++)
        { // attempt level i
            level[me] = i;
            victim[i] = me;
            // spin while conflicts exist
            for (int k = 0; k < n; k++)
                while ((k != me) && (level[k] >= i && victim[i] == me)){}
        }
    }
    /* This function sets the level of the thread to 0. */
    void unlock(int thread_id)
    {
        level[thread_id] = 0;
    }
};
```

- In constructor 2 arrays level and victim of size n are initialized
- When a thread acquires the lock
- For every level i:
 - Level[thread_id] is set to i
 - Victim[i] is set to thread_id
- while (($\exists k \neq me$) (level[k] >= i && victim[i] == me)) {}; is the spinning condition which can be translated to code as:
 - for (int k = 0; k < n; k++)
 - while ((k != me) && (level[k] >= i && victim[i] == me)){}
- For Unlocking simply level[thread_id] is set to 0

Tree-based Peterson Lock

Peterson Lock Node

A binary tree is created with $n/2$ leaves where each node is a slightly **modified Peterson lock** which takes in 2 threads for this lock those two threads are treated as 0,1

- When two threads enter a Peterson lock they are treated as 0,1 for that particular node
- Outside their thread_ids maybe anything
- To facilitate this mapping i have used two addition variables i_id,j_id

Internal 0th Thread_id= i_idth External thread_id

Internal 1th Thread_id = j_idth External thread_id

Each node has:

node *leftChild,*rightChild,*parent; => for binary tree

atomic_bool *flag , atomic_int victim; => As a simple peterson lock requires

int i_id,int j_id; => additional to facilitate n thread locking

Constructor

```
node() {}
node(node *par)
{
    flag = new atomic_bool[2];
    parent = par;
    i_id = -1;
    j_id = -1;
}
```

- Takes in par (parent Node) sets parent to par
- Initializes the flag array of atomic bool size 2
- Initializes i_id and j_id to -1

Lock and Unlock

- When a thread requests this Peterson lock first it checks if either `i_id`, `j_id` have never been set if that's the case then this incoming `thread_id` is mapped to 0,1 depending on the availability.
- If Both have been previously set it check if either the 0th or 1st thread in this Peterson lock is requesting the lock (`flag[i]` set to true),
 - if not then the incoming external id is mapped to that internal thread id is mapped
 - Else it waits for either thread to release the lock
- With this mapping in place, the Peterson lock works in the usual way.
- While Unlocking a `thread_id` simply the mapping is checked and corresponding 0th or 1st lock is release

```

void lock(int thread_id)
{
    //i_id is not set yet so set it to the thread_id requesting this lock
    if (i_id == -1)
        i_id = thread_id;
    // i_id is not set but j_id is not set so set j_id to the thread_id requesting
this lock
    else if (j_id == -1)
        j_id = thread_id;
    //both i_id and j_id are set
    else
    {
        //0th peterson thread is not requesting the lock
        if (flag[0].load() == false)
            i_id = thread_id;
        //1st peterson thread is not requesting the lock
        else if (flag[1].load() == false)
            j_id = thread_id;
        //both peterson threads are requesting the lock
        else
            //Wait till either frees up
            while (flag[0].load() == true && flag[1].load() == true){}
    }

    int i;
    if (thread_id == i_id)
        i = 0;
    else
        i = 1;

    int j = 1 - i;
    flag[i].store(true);
    victim.store(i);
    while (flag[j].load() && victim.load() == i){}
}

void unlock(int thread_id)
{
    int me;
    if (thread_id == i_id) me = 0;
    else if (thread_id == j_id) me = 1;
    flag[me].store(false);
}

```

Binary Tree of locks

Overview

- A binary tree is created with $n/2$ leaves where each node is a slightly **modified Peterson lock which takes in 2 threads for this lock those two threads are treated as 0,1**
- Each thread is assigned a leaf lock which it shares with one other thread.
- Each lock treats one thread as thread 0 and the other as thread 1.
- In the tree-lock's acquire method, the thread acquires every two-thread Peterson lock from that thread's leaf to the root.
- The tree-lock's release method for the tree-lock unlocks each of the 2-thread Peterson locks that thread has acquired.

Ex

- 16 thread we will create a binary tree with 8 leaf nodes and total $1+2+4+8 = 15$ nodes
- When thread 8 or 9 acquire the lock it acquires the leaf node- $4=8/2=9/2$ lock then makes its way to the root locking all the nodes in the path
- Similarly it release the lock it starts from the leaf node to the root node unlocking all the nodes in its path

0	Root							
1								
2								
3	0/1	2/3	4/5	6/7	8/9	10/11	12/13	14/15
	0	1	2	3	4	5	6	7

Constructor

- Leaf_nodes is a vector of leaf_nodes since other levels can be accessed using parents pointer i havent stored the other level vectors
- In constructor
 - For the 0th level i have created a vector and pushed the root to it
 - Then this vector is passed to the function recursiveBuild which builds the tree further

```
node *root; //root node of the tree

vector<node *> *leaf_nodes; //vector of leaf nodes

int height; //height of the tree

ptl(int n)

    this->n = n; //number of threads

    root = new node(NULL); //root node

    vector<node *> *tree = new vector<node *>; //initialize the tree

    tree->push_back(root); //push the root node to the tree

    leaf_nodes = recursiveBuild(tree); //recursively build the tree

    height = log2(n); //calculate the height of the tree
```

- This function takes in a vector of nodes
- Base case is the nodes vector size is n/2 ie that last level is reached.
- The level below the current nodes is built in new_nodes vector by iterating over nodes and creating and pushing their left and right Children in new_nodes

```
vector<node *> *recursiveBuild(vector<node *> *nodes)

{

    if (nodes->size() == ((this->n) / 2))

    {

        return nodes;

    }

    vector<node *> *new_nodes = new vector<node *>;
```

```

    for (int i = 0; i < nodes->size(); i++)
    {
        node *parent = nodes->at(i);
        node *left = new node(parent);
        node *right = new node(parent);
        parent->leftChild = left;
        parent->rightChild = right;
        new_nodes->push_back(left);
        new_nodes->push_back(right);
    }

    return recursiveBuild(new_nodes);
}

```

Lock and Unlock

- Locking and unlocking is easy we just need to traverse from leaf to root locking/unlocking the nodes in path

```

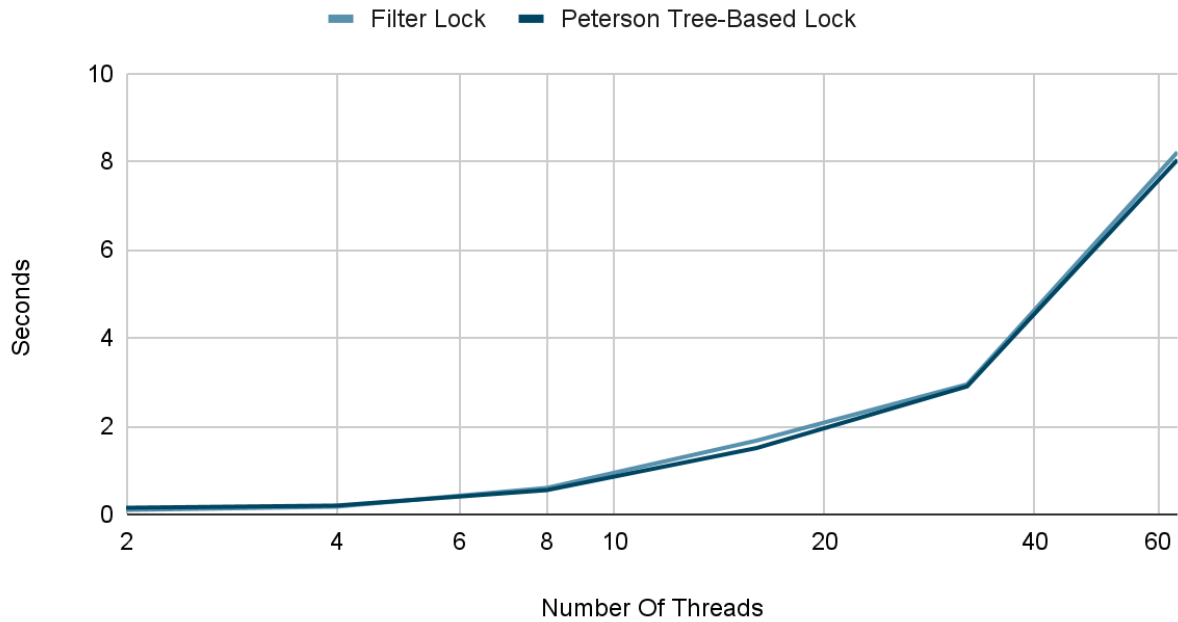
node *curr = (*leaf_nodes)[thread_id / 2];

while (curr)
{
    curr->(lock/unlock)(thread_id);
    curr = curr->parent;
}

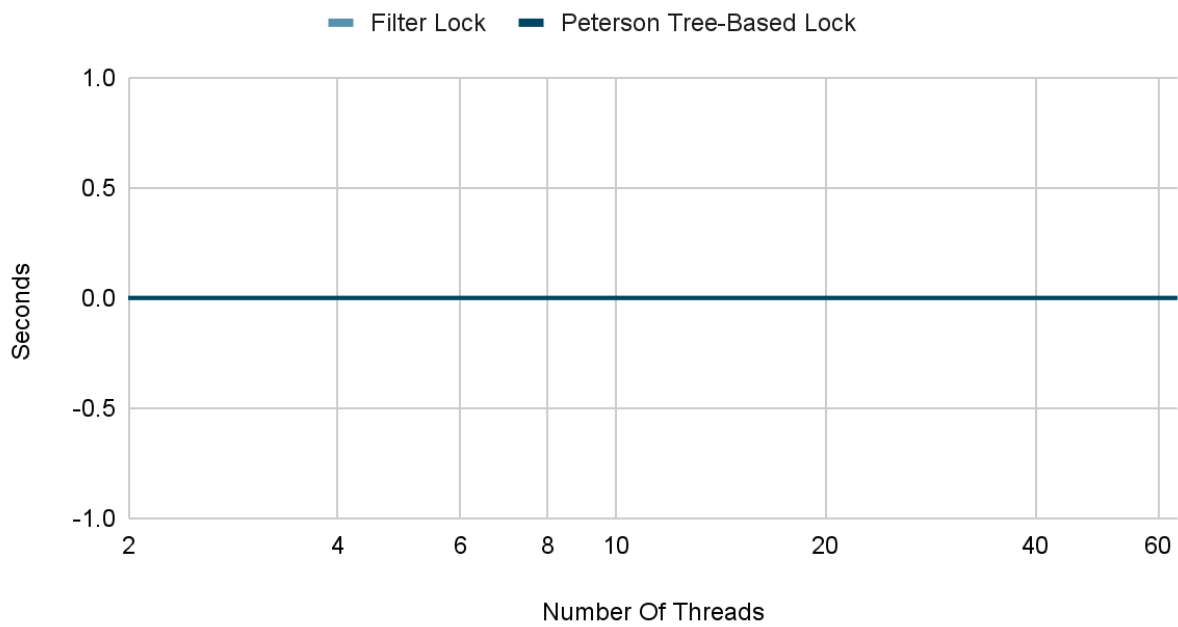
```


Analysis

CS Entry Time



CS Exit Time



CSEntry and CSExit Time= avg(avg(k CS by a thread))over n threads

When we vary the number of threads in 2^n we observe

- CSEntry Time: Increases exponentially with an exponential increase in the number of threads and both Filter Lock and Peterson Tree-based lock algorithms provide nearly similar performance with sometimes filter lock having an edge while other time tree lock gains an edge.

```
reqExitTime = getSysTime();
cout << i << "th CS Exit Reque
" << id << " (mesg 3)";
Test.unlock();
actExitTime = getSysTime();
```

- CSExitTime: Is almost zero in every case as CSExit time is the difference if we dive in further the difference comes out to be in a few microseconds as its the time taken in unlocking the lock which is a very less time-consuming task in Tree-based lock if the table grows very large then it may take some time as it will have to

unlock $\log_2(n)$ =height nodes

- On further trying out with 128 threads some Exit time we can see in seconds.

```
9th CS Exit Request At 2022-02-12 13:27:45:740 by thread 87 (mesg3)
9th CS Entry At 2022-02-12 13:27:45:740 by thread 79 (mesg2)
9th CS Exit Request At 2022-02-12 13:27:45:799 by thread 79 (mesg3)
9th CS Entry At 2022-02-12 13:27:45:799 by thread 95 (mesg2)
9th CS Exit Request At 2022-02-12 13:27:45:850 by thread 95 (mesg3)

-----
Average CS Entry Time for Filter Lock: 18.7852
Average CS Exit Time for Filter Lock: 0.003125
Average CS Entry Time for PTL: 20.7328
Average CS Exit Time for PTL: 0.0015625
-----
./a.out 2641.35s user 23.69s system 645% cpu 6:52.94 total
(base) * [master] [~/Data/Semester/Assignments/ProgAssn2-CS19B1017]$
```

References:

- The Art of Multiprocessor Programming Maurice Herlihy, Nir Shavit